

CS302 -- Lab 4 -- Enumeration

- CS302 -- Data Structures and Algorithms II
- [James S. Plank](#)
- [This file: http://web.eecs.utk.edu/~jplank/plank/classes/cs302/Labs/Lab4](http://web.eecs.utk.edu/~jplank/plank/classes/cs302/Labs/Lab4)
- Lab Directory: /home/jplank/cs302/Labs/Lab4

What you hand in

You will submit the single file `matrix_enum.cpp`. The TA's will compile that to an executable `matrix_enum` and test it with the grading script.

Introduction, and your enumerations

Back in the late 2000's and early 2010's, I was doing research on *erasure codes*. These are techniques for protecting data in storage systems composed of multiple disks (think [RAID](#)). In 2008, I wrote a paper entitled "A New Minimum Density RAID-6 Code with a Word Size of Eight." If you dare, you can read the paper at <http://web.eecs.utk.edu/~jplank/plank/papers/NCA-2008.html>. Read it out loud to your significant other when he/she can't sleep.

Part of this research was performing enumerations of square matrices. The elements in the matrices can be any of three characters: '.', 'x' and 'E'. The enumeration is parameterized by two values:

- **W** (which stands for the "word size") is the number of rows in the matrices. It is also the number of columns, because the matrices are square.
- **E** (which stands for "extra non-zero entries") is a number between 0 and $W^2 - W$.

Now, given **W** and **E**, your job is to enumerate all matrices of the following form:

- There are exactly **W** elements with x's. There must be an 'x' in every row and every column. (BTW, this is called a [Permutation Matrix](#)).
- There are exactly **E** elements that have E's in them, and they can't be where the x's are.
- The remaining $W^2 - W - E$ elements contain '.'.

For example, when **W** is 3 and **E** is 1, then you have the following 36 matrices:

XE.	X.E	X..	X..	X..	X..	XE.	X.E	X..	X..	X..	X..
.X.	.X.	EX.	.XE	.X.	.X.	.X.	.X.	E.X	.EX	.X.	.X.
..X	..X	..X	..X	E.X	.EX	.X.	.X.	.X.	.X.	EX.	.XE
EX.	.XE	.X.	.X.	.X.	.X.	EX.	.XE	.X.	.X.	.X.	.X.
X..	X..	XE.	X.E	X..	X..	.X.	.X.	E.X	.EX	.X.	.X.
..X	..X	..X	..X	E.X	.EX	X..	X..	X..	X..	XE.	X.E
E.X	.EX	..X	..X	..X	..X	E.X	.EX	..X	..X	..X	..X
.X.	.X.	EX.	.XE	.X.	.X.	X..	X..	XE.	X.E	X..	X..
X..	X..	X..	X..	XE.	X.E	.X.	.X.	.X.	.X.	EX.	.XE

Your job is to write a program called `matrix_enum`, which takes three command line arguments: **W**, **E** and either an 'x' or an 'h'. Your program will enumerate all of the matrices for **W** and **E**, in any order you want, and print them out in one of the following two formats:

- If the last argument is 'x', then you'll print the matrices out in the format above. You should print each matrix as **W** lines of **W** characters that are '.', 'x' or 'E'. After each matrix, you print a blank line.
- If the last argument is 'h', then you'll convert each line of each matrix into an integer, and print that integer in hexadecimal, with no leading 0's and no leading "0x". If element *i* in a row is 'x' or 'E', then you'll set the *i*-th bit of the number to 1. Otherwise, the *i*-th bit is zero. You'll print each integer on its own line, and print a blank line at the end of each matrix.

Here are a few runs:

UNIX> <code>matrix_enum 2 0 x</code> X. .X .X X. UNIX>	UNIX> <code>matrix_enum 2 1 x</code> XE .X X. EX EX X. .X XE UNIX>	UNIX> <code>matrix_enum 2 2 x</code> XE XE EX XE UNIX>	UNIX> <code>matrix_enum 2 0 h</code> 1 2 2 1 UNIX>	UNIX> <code>matrix_enum 2 1 h</code> 3 3 2 1 3 3 1 2 3 UNIX>	UNIX> <code>matrix_enu</code> 3 3 3 3 3 UNIX>
---	---	---	---	--	---

Approach

This is a two-level enumeration. The first enumeration is enumerating the permutation matrices (the x's). You can do this by enumerating permutations of the numbers (0,1,2,...**W**-1). Let's suppose that each permutation is represented by a vector of **W** numbers. Then, the 'x' in row *i* will be in the column specified by element *i* of the vector. For example, if **W** equals 3, then the permutation (1,2,0) represents the matrix:

.X.
..X
X..

I don't want you to use `next_permutation()` to implement this part. Use the recursive technique specified in the [Enumeration lecture notes](#), in the section entitled "Using Recursion to Permute." If you use `next_permutation()`, you will lose 10 points on this lab, so when you're developing your code, you may want to start by using `next_permutation()`, and then changing it to use recursion when you've gotten everything else working.

At the second level, you have $W^2 \cdot W$ potential locations for the 'E' characters. Thus, this is an " n choose k " where n is equal to $W^2 \cdot W$ and k is equal to E . Implement this using the recursive technique specified in the [Enumeration lecture notes](#), in the section entitled "Using Recursion to Perform an " n choose k " Enumeration."

When I did this, I created a vector out of the non-'x' matrix elements. I stored these as numbers: (row * W + column). I then enumerated all ways to choose E of these, using the same technique as in the lecture notes. Continuing the example above, suppose W equals three, and my permutation matrix is (1,2,0). Then the potential spots for 'E' elements are the following (again, using (row * W + column)):

(0, 2, 3, 4, 7, 8).

Here is my main class definition. You don't need to use this, but you may find it useful:

```
class Matrix {
public:
    int W;
    int E;
    int P;
    vector<int> Perm;          /* This is 'x' or 'h' */
    vector<int> Non_X;        /* Permutation of 0 .. (W-1), for the 'X' elements. */
    vector<int> E_ID;         /* This is the row/col id of each of the non-X elements. */
    void Print();             /* This is the row/col id of the E elements */
    void Permute(int index);  /* Print the matrix defined by W, Perm and E_ID */
    void Choose(int index);   /* This is the recursive permuting method. */
};
```

The Gradescript

I have two techniques for grading these. The first pipes the output of your program through the program [flatten_for_grading.cpp](#). This program coalesces the lines for each matrix down to one line, and then prints the lines. Thus, if you pipe the output for this to `sort`, then you'll get all of the matrices, printed on one line, sorted lexicographically

```
UNIX> matrix_enum 2 1 x | flatten_for_grading | sort
.X XE
EX X.
X. EX
XE .X
UNIX> matrix_enum 2 1 h | flatten_for_grading | sort
1 3
2 3
3 1
3 2
UNIX>
```

Your output, after being piped through `flatten_for_grading` and then `sort`, has to match my output (piped through `flatten_for_grading` and `sort`) verbatim. This will be the technique used for gradescripts 0 through 70.

The second technique is to use the program `double_check`. This takes W , E , and $h|x$ on the command line and then accepts input of your program piped through `flatten_for_grading`. It then checks every line of input to see if it is legal for those parameters. This means that for 'x':

- Lines have the proper number of characters and are in the proper format.
- Each row has exactly one 'x' in a different column.
- The number of 'E' and '.' characters is correct.
- Lines are unique.

For 'h', the line is converted into a string, and then checked as in 'x' above. There is a little subtlety here, because a single 'h' line can correspond to quite a few strings. What `double_check` does is enumerate all of the strings that the 'h' line corresponds to, and it counts the legal ones. It then allows to use that string that many times. Let's try an example. Suppose that W equals three and E equals 2, and you specify the 'h' matrix (3,3,4). This can correspond to the following legal matrices:

XE.	EX.
EX.	XE.
. .X	. .X

So `double_check` will allow you to specify (3,3,4) twice, but not three times.

The `double_check` program prints nothing if its input is legal. If you have a problem, it prints an error message on standard output. For example, here's what happens if you specify (3,3,4) three times in the above example:

```
UNIX> double_check 3 2 h
3 3 4
3 3 4
3 3 4
Bad line 3: Too many lines (3) with these values.
Here are the matrices that correspond to these values:
EX. XE. . .X
XE. EX. . .X
UNIX>
```

Gradescripts 71 through 100 pipe the output of your program through `flatten_for_grading`, and then through `head -n 50000`, and then through `double_check`. This lets me grade for larger values of W and E without waiting for the giant enumerations to complete.

To help you manage your expectations -- the following table contains how long the gradescripts took on my programs (seconds). Most of this time is in `double_check`, by the way -- please see [this Piazza question and answer for more detail on why gradescript 81 takes so long](#).

1 - 0.08	21 - 3.31	41 - 0.06	61 - 0.29	81 - 11.78
2 - 0.06	22 - 0.76	42 - 2.71	62 - 0.04	82 - 1.00
3 - 0.06	23 - 0.43	43 - 0.13	63 - 0.29	83 - 0.77
4 - 0.06	24 - 0.09	44 - 2.01	64 - 0.35	84 - 1.53
5 - 0.06	25 - 1.38	45 - 0.10	65 - 0.25	85 - 1.52
6 - 0.06	26 - 0.32	46 - 0.06	66 - 0.68	86 - 2.20
7 - 0.07	27 - 1.71	47 - 1.27	67 - 0.06	87 - 0.53
8 - 0.07	28 - 1.81	48 - 0.15	68 - 0.50	88 - 0.52
9 - 0.06	29 - 0.14	49 - 0.22	69 - 0.27	89 - 0.41
10 - 0.06	30 - 3.17	50 - 0.18	70 - 0.22	90 - 2.65
11 - 1.65	31 - 0.34	51 - 0.04	71 - 0.27	91 - 2.40
12 - 0.06	32 - 0.83	52 - 0.04	72 - 1.73	92 - 1.48
13 - 0.81	33 - 0.44	53 - 0.04	73 - 0.42	93 - 1.80
14 - 0.20	34 - 2.77	54 - 0.04	74 - 0.05	94 - 0.78
15 - 4.09	35 - 1.54	55 - 0.40	75 - 1.21	95 - 0.63
16 - 0.37	36 - 0.07	56 - 0.05	76 - 1.17	96 - 1.76
17 - 0.09	37 - 0.07	57 - 0.05	77 - 1.68	97 - 1.17
18 - 1.93	38 - 0.39	58 - 0.05	78 - 1.79	98 - 0.57
19 - 0.66	39 - 2.07	59 - 0.04	79 - 0.49	99 - 0.30
20 - 0.21	40 - 3.38	60 - 0.04	80 - 2.08	100 - 2.38